

Integration of Hardware–BCP for SAT Acceleration

EECS 219C Final Report

Evan Wong

Ansh Shah

Spring 2026

Abstract

SAT solvers underpin every modern formal verification flow, but Boolean Constraint Propagation (BCP) is well known to consume the overwhelming majority of CDCL runtime.¹ We design a hardware–software coprocessor that pushes BCP into a CAM/SRAM RTL engine while keeping CDCL search logic (VSIDS, 1–UIP, non–chronological backtracking) in a minimal C++ wrapper, simulated through Verilator. On SATLIB Uniform Random 3–SAT phase–transition benchmarks, the coprocessor delivers **37–341×** faster BCP than Kissat 4.0.4, outperforms Kissat, MiniSat2, and Glucose end–to–end up to ~ 175 variables, and reveals that beyond that point the bottleneck shifts to the CDCL heuristic — confirmed by a counter–factual that re–runs the search with software BCP rates.

1 Introduction & Problem Definition

SAT solving is the engine behind industrial formal verification: Cadence JasperGold, Synopsys VC Formal, and Siemens Questa Formal all rely on CDCL solvers for bounded model checking and property checking. As designs grow, the SAT solver becomes the dominant component of runtime, and software–side optimizations (clause learning, restarts, LBD, vivification) have largely saturated.

Profiling the canonical CDCL loop shows a striking distribution (Figure 1): **BCP alone is 82.7%** of total solver time, with conflict analysis (8.7%), literal redundancy (3.9%), and branching heuristics (0.6%) trailing far behind — consistent with the empirical bottleneck identified in the course material. This makes BCP the natural target for hardware acceleration — accelerating anything else would deliver, by Amdahl’s Law, at most a few percent.




Figure 1: Average runtime distribution across CDCL operations on phase–transition benchmarks. BCP dominates by an order of magnitude.

¹S. A. Seshia, EECS 219C lecture notes, “Boolean Satisfiability Solving.” <https://people.eecs.berkeley.edu/~sseshia/219c/lectures/SATsolving.pdf>

Problem statement. Design and verify a hardware BCP engine that delivers single-cycle unit propagation, integrate it with a minimal CDCL software wrapper, and measure end-to-end speedup against the most competitive software solvers (Kissat 4.0.4, MiniSat2, Glucose) on standardized benchmarks.

2 Approach and Related Work

Why CAM + SRAM. Prior hardware SAT work² consistently identifies *memory bandwidth*, not arithmetic, as the limiter. Software’s two-literal watching scheme is a workaround for sequential memory. In hardware, a CAM collapses “find every clause containing v ” to a single associative lookup, and an SRAM holds the per-variable assignment used by combinational clause-evaluation logic — so every clause is re-evaluated every cycle in parallel and watched literals are unnecessary. We adopt the HW-BCP algorithm and CAM/SRAM cell behavior of Park et al. (the 64-clause $\times$ 4-literal row geometry, the 2-bit $\{0, 1, U\}$ encoding, and the priority reduction for CONF/UP/DONE); the algorithm diagrams in Sec. 3 are reproduced from that work. The integration into a CDCL coprocessor, the Verilator harness, and the timing methodology are ours.

Where we differ. Earlier proposals are typically all-hardware (HW also performs branching), which makes them inflexible and forces the CDCL heuristic into silicon. Our design is a **coprocessor**: HW handles only BCP; SW handles every heuristic — VSIDS, 1-UIP analysis, restarts, learnt management. This isolates the BCP speedup so it can be measured cleanly, and keeps the door open for arbitrary CDCL improvements without an RTL respin.

Architectural sketch. Figure 2 shows the target geometry. The chip stores up to 1024 clauses (capacity exhausted in our 1024-slot RTL prototype; a 1M-slot SRAM is feasible in TSMC 65 nm using Park et al.’s custom memory cells, and is what we model in our timing extrapolation). Each clause is a row of up to 4 literals; a CAM cell holds the variable ID and an SRAM cell holds the 2-bit value (0, 1, or U = unassigned). A single combinational reduction returns three flags per clause every cycle: CONFLICT, UNIT, and DONE.




Figure 2: Coprocessor partitioning. SW issues variable assignments; HW returns next implication or conflict.

²Y. Zhao & K. A. Sakallah, ICCAD 2008; J. D. Davis et al., DAC 2008; S. Park et al., <https://par.nsf.gov/servlets/purl/10215919>.

3 Formal Model and Algorithms

3.1 Clause and trail semantics

A CNF formula is $F = \bigwedge_{c \in C} c$, where each clause is a disjunction of literals $c = \ell_1 \vee \dots \vee \ell_k$, $k \leq 4$. A literal is $\ell = (var, pol)$ with $pol \in \{0, 1\}$. The solver maintains a partial assignment $\sigma : V \rightarrow \{0, 1, U\}$ on a *trail* (decision stack with implication reasons).

Per-clause evaluation. Define for each literal ℓ under assignment σ :

$$lv(\ell, \sigma) = \begin{cases} 1 & \sigma(\text{var}(\ell)) = \text{pol}(\ell) \\ 0 & \sigma(\text{var}(\ell)) \neq \text{pol}(\ell), \neq U \\ U & \sigma(\text{var}(\ell)) = U \end{cases}$$

A clause c is then classified as

$$\text{state}(c) = \begin{cases} \text{DONE} & \exists \ell \in c. lv(\ell) = 1 \\ \text{CONFLICT} & \forall \ell \in c. lv(\ell) = 0 \\ \text{UNIT} & \exists! \ell \in c. lv(\ell) = U \wedge \text{rest are 0} \\ \text{PENDING} & \text{otherwise} \end{cases}$$

The combining logic produces a global flag with priority $\text{CONFLICT} \succ \text{UNIT} \succ \text{DONE}$.

3.2 HW-BCP cycle

Each clock cycle, the engine executes one of four operations (Table 1); BCP is the only one that performs a parallel reduction over all N clauses.

Op	Description	Cycles
LOAD	Write one clause literal into a CAM row	1
BCP	Return next unit/conflict from N rows	1
UNDO	Clear a variable assignment	1
SW_BCP	Overflow learnt propagated in SW (HW-counted) [†]	1

Table 1: Hardware operations. [†]Modeled as 1 HW cycle to reflect a production 1M-slot SRAM in which all clauses live in HW.

At a target frequency of ~ 99 MHz (a 10.1 ns critical path, measured from synthesis with parameterized memories), the total hardware time is

$$T_{\text{hw}} = (n_{\text{load}} + n_{\text{bcp}} + n_{\text{undo}} + n_{\text{sw_bcp}}) \times 10.1 \text{ ns.}$$

3.3 Coprocessor algorithm

The CDCL outer loop is essentially MiniSat-style 1-UIP, but every `propagate()` call traps into the RTL (Algorithm 1).

Algorithm 1 HW/SW Coprocessed CDCL

```
1: Load all original clauses via LOAD
2: while true do
3:    $r \leftarrow \text{HW-BCP}()$  (single-cycle parallel scan)
4:   if  $r = \text{CONFLICT}$  then
5:     if decision level = 0 then
6:       return UNSAT
7:     end if
8:      $(c_{\text{learnt}}, lvl) \leftarrow \text{Analyze1UIP}()$ 
9:     Backtrack( $lvl$ ) via UNDO
10:    LoadLearnt( $c_{\text{learnt}}$ ) (or overflow  $\rightarrow$  SW pool)
11:  else if  $r = \text{UNIT}(\ell)$  then
12:    Enqueue  $\ell$  on trail; LOAD updates SRAM
13:  else
14:    if trail covers all vars then
15:      return SAT
16:    end if
17:     $v \leftarrow \text{VSIDS-Pick}()$ ; assign and LOAD
18:  end if
19: end while
```

3.4 Worked example

Figure 3 shows a single propagation step for $C_0 = x_1 \vee x_2 \vee x_3$. The trail head is $x_2 = 0$. The CAM returns rows referencing x_1, x_2, x_3 for C_0 ; the SRAM resolves $x_1 = 0, x_2 = 0, x_3 = U$; the per-clause evaluator detects exactly one unassigned literal with the rest false, so it asserts UNIT and emits $x_3 \leftarrow 1$ in the same cycle.

var	cls	idx
x_1	C_0	0
x_2	C_0	1
x_1	C_1	0
x_3	C_1	1

var	val
x_1	0
x_2	0
x_3	U
x_4	U

$C_0 = x_1 \vee x_2 \vee x_3$
 $x_1 = 0$ ✗
 $x_2 = 0$ ✗
 $x_3 = U \Rightarrow$
UNIT: $x_3 \leftarrow 1$

Figure 3: Single-cycle BCP: CAM returns matching rows, SRAM gives literal values, evaluator detects UNIT.

4 Major Technical Challenges

C1. Time bookkeeping in a Verilator harness. Wall time conflates real CDCL, Verilator `eval()` cost, and simulation-only $O(N)$ scans that would be $O(1)$ in silicon. We accumulate the latter two into hw_sim_ns and hw_aux_ns and subtract: $T_{sw} = T_{total} - \Delta hw_sim_ns - \Delta hw_aux_ns$. An early bug had the timers running from solver construction, so initial clause loads polluted the BCP budget; we now snapshot at `solve()` entry.

C2. CAM capacity vs. learnt blowup. On uf200–860 the originals fill 84% of the 1024-slot CAM, leaving ~ 164 slots before overflow. We handle this with (a) a SW overflow pool whose cycles

are attributed to HW (justified by the 1M-slot production target), and (b) a hard guard that aborts when blowup exceeds capacity (Figure 4).




Figure 4: CAM slot utilization on each benchmark set. uf200 leaves only ~ 164 free slots, causing immediate learnt overflow and degraded BCP effectiveness.

C3. HW/SW state coherence on backtrack. A bug in `undo_hw_/load_hw_` caused the SRAM polarity bit to lag the CAM by one cycle under rapid load/undo bursts, producing spurious unit implications. The fix drives `val_we` synchronously with the load address and gates read-back behind one extra cycle of stable state.

C4. Full-hierarchy elaboration timeout. A pure-RTL run on the full $1024 \times 8 \times 64 = 524,288$ -clause hierarchy timed out in Verilator (behavioral memories explode the logic cone). We therefore parameterize the RTL so benchmarks elaborate a single 1024-slot `sat_submodule`, while the full `sat_array` still elaborates and lints to prove structural completeness.

5 Results

We benchmarked on SATLIB Uniform Random 3-SAT phase-transition instances ($c/v \approx 4.27$) with up to 200 variables and 860 clauses. Each set contains 100–1000 instances; we report median times. All times are reported in μs , recomputed at the synthesis-profiled 10.1 ns per HW cycle.

5.1 End-to-end median solve time

What worked. The coprocessor is faster than *every* software solver on every instance set with at most ~ 150 variables, often by 5–30 $\times$. The mechanism is exactly what the architecture promised: BCP per propagation drops from hundreds of nanoseconds (Kissat) to a single 10.1 ns cycle.

What didn't. On ≥ 175 variables, Kissat overtakes us on SAT instances. The cause is not BCP throughput (Section 5.3) but the lack of restarts, phase saving, and clause vivification in our minimal CDCL wrapper: once the search tree gets deep, raw BCP speed cannot make up for suboptimal decisions.




Figure 5: Median solve time on UF/UUF benchmarks, log-scaled. Our coprocessor dominates up to ~ 175 variables; beyond that, Kissat’s superior CDCL heuristic surpasses our naive search.

Vars	SAT wins (vs. Kissat)	UNSAT wins	Gap
75	100/100	100/100	–
100	959/1000	1000/1000	+4 pp
125	86/100	100/100	+14 pp
150	71/100	98/100	+27 pp
175	60/100	93/100	+33 pp

Table 2: Win rate vs. Kissat by satisfiability. UNSAT instances favor our coprocessor far more than SAT instances do. We attribute this to Kissat’s aggressive phase saving / restarts being most useful on satisfiable inputs — a heuristic edge that does not help when the proof must enumerate the whole tree.

5.2 SAT vs. UNSAT asymmetry

5.3 Counterfactual: is BCP really accelerated?

To attribute the speedup to HW-BCP (rather than a Verilator quirk), we re-ran the exact CDCL trajectory but plugged in Kissat’s measured BCP rate per propagation. Our solver runs $2.2\text{--}15.7\times$ slower in this counterfactual (Figure 6) — HW-BCP is load-bearing, even on UF175 where Kissat wins overall: the CDCL heuristic, not BCP, is the remaining bottleneck.

5.4 Amdahl’s-law breakdown

Our solver spends $\sim 93\%$ of its time in software CDCL and $\sim 7\%$ in BCP — the inverse of MiniSat2, where BCP is $\sim 75\%$ of runtime. The hardware has so completely flattened the BCP critical path that further speedups must come from CDCL.

5.5 Small but provably hard instances

Phase-transition random-3-SAT is *statistically* hard; worst-case CDCL is set by adversarial classes admitting only super-polynomial resolution proofs. We built a small suite at `tests/satlib/hard_instances/`




Figure 6: Replacing HW-BCP with Kissat’s measured BCP rate, while keeping the identical CDCL trajectory, slows the same solver by 2.2–15.7 $\times$. The hardware is doing the work.




Figure 7: Total runtime breakdown on uf100–430 (median). Our solver is now CDCL-bound; MiniSat2 is BCP-bound.

(clauses wider than four literals are Tseitin-split to fit the 4-lit RTL): pigeonhole PHP_n ($n=3..8$, the canonical Haken exponential-resolution lower bound); Tseitin parity on odd cycles (C_5, C_7, C_9); and k -colourability of the Grttsch graph ($\chi=4$, so $k=3$ is hard UNSAT and $k=4$ is SAT).

Methodology. Kissat/MiniSat2 self-report CPU time at ~ 10 ms granularity, flooring sub-ms runs to zero; we fall back to wall-clock minus a per-binary subprocess-startup baseline (≈ 4 – 5 ms) when the self-clock is 0. HW-BCP figure is $hw_cycles \times 10.1 \text{ ns} + sw_time$ as elsewhere.

Discussion. On PHP, every solver must enumerate the same exponentially many conflicts; the win is decided by *cost-per-propagation*. SW two-literal watching degrades on dense binary clause bases (PHP is mostly width-2) while HW parallel re-evaluation pays no such penalty — the same effect behind the 30–90 $\times$ wins on Tseitin and Grttsch. $\text{PHP}_3/\text{PHP}_5$ solve below the wall-clock floor for some SW solvers, so we make no claim there. PHP_9 and beyond would push us into the overflow pool and erode the constant-cycle BCP, which the 1M-slot production target eliminates.

5.6 Future work

- Replace the MiniSat-style heuristic with Kissat-style restarts, phase saving, and LBD-based clause deletion. Closing the CDCL gap should restore HW dominance past 175 variables.
- Push the RTL through TSMC 65 nm using Park et al.’s custom CAM/SRAM cells (the only PDK for which those cells were characterized); a 1M-slot SRAM in that node eliminates the overflow pool entirely.

Instance	Exp.	HW–BCP (μ s)	Kissat (μ s)	MiniSat2 (μ s)	Glucose (μ s)	Win
PHP ₃	UNSAT	66	<1k	1 106	1 127	noise
PHP ₄	UNSAT	184	684	2 710	1 005	HW
PHP ₅	UNSAT	731	3 891	<1k	2 665	noise
PHP ₆	UNSAT	3 352	10 000	7 911	5 930	HW
PHP ₇	UNSAT	6 034	10 000	37 292	61 267	HW
PHP ₈	UNSAT	6 803	40 000	190 585	962 534	HW
Tseitin–C ₅	UNSAT	12	1 419	1 357	920	HW
Tseitin–C ₇	UNSAT	19	947	2 254	828	HW
Tseitin–C ₉	UNSAT	21	895	962	1 458	HW
Grttsch–3col	UNSAT	87	550	1 404	1 167	HW
Grttsch–4col	SAT	19	468	1 394	1 170	HW

Table 3: Hard-instance suite, single-run. “noise” = result falls below subprocess-startup floor for at least one solver, so no winner is claimed. HW wins outright on 9 of 11 instances.




Figure 8: Pigeonhole PHP_n scaling, $n = 3$ –8, log-scale. The gap between HW–BCP and the SW solvers *widens* as n grows: on PHP₈, HW is 5.9 $\times$ faster than Kissat, 28 $\times$ faster than MiniSat2, and 142 $\times$ faster than Glucose — on a class where the search space itself is exponential, the constant-time HW–BCP cycle compounds across every conflict in the proof.

- Pipeline LOAD and BCP so a decision and the following propagation overlap, halving CDCL–HW round-trip latency.

6 Team Roles

Evan Wong primarily took charge in collaborating with staff such as Pierluigi Nuzzo and course staff from EECS 151 for RTL infrastructure and the RTL: the behavioral CAM/SRAM array (`cam_sram_array.sv`, `cam_cell.sv`, `sram_row.sv`), the per-clause evaluator (`processing_logic.sv`), the FSM driving load/BCP/undo, and the Verilator integration; he also wrote the testbench harness at the RTL level. **Ansh Shah** owned the CDCL software: VSIDS scoring, 1–UIP conflict analysis, learnt management with the SW overflow pool, and the benchmark/parsing infrastructure (`bench_cdcl.py`) and others across multiple iterations. The counterfactual analysis (Section 5.3) and the SAT/UNSAT asymmetry study were a joint effort, as was the slide deck and this report.

7 Course Topics Applied

Almost every layer of this project draws on course material:

- **The CDCL loop, unit rule, and BCP as the dominant cost.** The coprocessor implements the canonical CDCL algorithm (preprocess, decide, deduce, analyze, backtrack); the unit rule is the UNIT flag emitted by the per-clause evaluator. The class’s empirical claim that BCP dominates runtime is the entire reason we built a HW-BCP coprocessor, and Figure 7 is the post-hoc confirmation.
- **Memory accesses per step.** The lecture’s principle of minimizing memory traffic per propagation underlies SW two-literal watching; our HW achieves the same end via per-cycle parallel evaluation, making cost $O(1)$ in clause width and watched literals unnecessary.
- **VSIDS, 1-UIP, non-chronological backtracking.** Our SW wrapper uses VSIDS scoring with periodic decay, performs 1-UIP conflict analysis over the implication graph, and jumps to the asserting level via UNDO cycles — all directly from class.
- **Phase transitions and resolution proofs.** SATLIB UF/UUF at $c/v \approx 4.27$ sits on the empirical 3-SAT hardness peak. The hard-instance suite (Section 5.5) leverages the resolution-proof framing: any CDCL refutation of pigeonhole encodes a resolution proof, and pigeonhole has only super-polynomial ones — so conflict count is invariant across solvers and HW wins on cost-per-propagation.
- **Clause deletion / learnt management.** Our SW overflow pool plus the 1024-slot CAM cap is the resource-constrained version of the standard clause-deletion problem covered in class.

8 Feedback

What helped. The class’s depth on CDCL internals (1-UIP analysis, VSIDS bumping, restart strategy, LBD) mapped directly onto our SW wrapper. The discussion of watched literals was critical because understanding *why* software needs them clarified *why* hardware can skip them. The phase-transition discussion gave us our benchmark.

What we wish had been covered. A short walkthrough of modern CDCL engineering choices beyond the textbook 1-UIP/VSIDS core — restart policies, phase saving, LBD-based clause deletion, and clause vivification as used in Kissat — would have helped us read the heuristic gap our counterfactual exposed and pick which upgrade to prototype next. Otherwise the class material covered what the project needed; the remaining gaps were hardware-specific and reasonably out of scope.

Code Availability

Full source (RTL, C++ coprocessor, benchmark scripts, hard-instance generator and CNFs): <https://github.com/The-Ansh-Shah/bcp-hw-accelerator>.